

# MongoDB 3.2 – Document Validation

Andrew Morgan  
@andrewmorgan |  
andrew.morgan@mongodb.com

29<sup>th</sup> October 2015



mongoDB

DISCLAIMER: MongoDB's product plans are for informational purposes only. MongoDB's plans may change and you should not rely on them for delivery of a specific feature at a specific time.

# Agenda

Value of flexible schemas

Downside of flexible schemas

What 3.2 adds

What 3.2 doesn't add

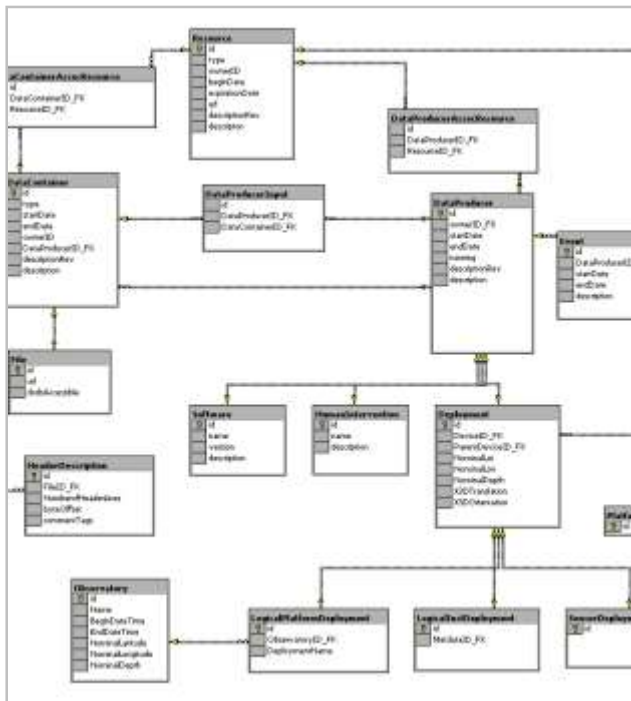
Options

Production lifecycle to add Document Validation

Walkthrough

# Power of flexible schemas

# RDBMS



# MongoDB

```
{
  _id: ObjectId("4c4ba5e5e8aabf3"),
  employee_name: {First: "Billy",
                  Last: "Fish"},
  department: "Engineering",
  title: "Aquarium design",
  pay_band: "C",
  benefits: [
    { type: "Health",
      plan: "PPO Plus" },
    { type: "Dental",
      plan: "Standard" }
  ]
}
```

# Power of flexible schemas

```
{
  _id: ObjectId("4c4ba5e5e8aabf3"),
  employee_name: {
    First: "Billy",
    Last: "Fish"},
  department: "Engineering",
  title: "Aquarium design",
  pay_band: "C",
  benefits: [
    {type: "Health",
      plan: "PPO Plus" },
    {type: "Dental",
      plan: "Standard" }
  ]
}
```

- Relational
  - Up-front schema definition phase
  - Adding new column takes time to develop & \*lots\* of time to roll out in production
    - Existing rows must be reformatted
- MongoDB:
  - Start hacking your app right away
  - Want to store new type of information?
    - Just start adding it
    - If it doesn't apply to all instances – just leave it out

# Power of flexible schemas

```
{
  _id: ObjectId("4c4ba5e5e8aabf3"),
  employee_name: {
    First: "Billy",
    Last: "Fish"},
  department: "Engineering",
  title: "Aquarium design",
  pay_band: "C",
  benefits: [
    {type: "Health",
      plan: "PPO Plus" },
    {type: "Dental",
      plan: "Standard" }
  ]
}
```

- Relational
  - Up-front schema definition phase
  - Adding new column takes time to develop & \*lots\* of time to roll out in production
    - Existing rows must be reformatted
- MongoDB:
  - Start hacking your app right away
  - Want to store new type of information?
    - Just start adding it
    - If it doesn't apply to all instances – just leave it out

# Why validate documents?


- Many people writing to the database
  - Many developers
  - Many teams
  - Many companies
  - Many development languages
- Multiple applications want to exploit the same data, need to agree on what's there
- Usually a core subset of keys you always want to be there
- For any key may care about:
  - Existence
  - Type
  - Format
  - Value
  - Existence in combination with other keys (e.g. need a phone number **or** an email address)

# Why validate documents?

- Good to have a 'contract' for what's in a collection
  - When reading from the “subscriber” collection, I know that every document will include a subscription plan name:

```
db.subscriptions.find(  
  name: "Billy Fish",  
  $and: [{plan: {$exists: true}},  
         {plan: {$type: 2}}])
```

- < MongoDB 3.2, this is an application responsibility
  - 3rd party tools like Mongoose can help
- Best implemented as a layer between the application and driver (Data Access Layer)





**Get the database to do the work!**

# Document Validation - MongoDB 3.2

- Configure Document **Validation within the database**
- Use familiar **MongoDB Query Language**
- Automatically tests each insert/update; delivers **warning or error** if a rule is broken
- You choose **what keys to validate and how**

```
db.runCommand({
  collMod: "contacts",
  validator: {
    $and: [
      {year_of_birth: {$lte: 1994}},
      {$or: [
        {phone: { $type: "string"}},
        {email: { $type: "string"}}
      ]}]
  }
})
```

# Document Validation - MongoDB 3.2

```
db.getCollectionInfos({name:"contacts"})
[
  {
    "name": "contacts",
    "options": {
      "validator": {
        "$and": [
          {"year_of_birth": {
            "$lte": 1994}},
          {"$or": [
            {"phone": {"$type": "string"}},
            {"email": {"$type": "string"}}
          ]}
        ]},
      "validationLevel": "strict",
      "validationAction": "error"
    }
  }
]
```

# Document Validation - MongoDB 3.2

```
db.contacts.insert(  
  name: "Fred",  
  email: "fred@clusterdb.com",  
  year_of_birth: 2012  
))
```

## Document failed validation

```
WriteResult({  
  "nInserted": 0,  
  "writeError": {  
    "code": 121,  
    "errmsg": "Document failed validation"}})
```

# Document Validation - MongoDB 3.2

```
db.runCommand({collMod: "bleh",  
              validator: {  
                rogue: {$exists:false}  
              }  
            }  
          });
```

# Document Validation - MongoDB 3.2

- Can check most things that work with a `find` expression
  - Existence
  - Non-existence
  - Data type of values
  - `<`, `<=`, `>`, `>=`, `==`, `!=`
  - AND, OR
  - Regular expressions
  - Some geospatial operators (e.g. `$geoWithin` & `$geoIntersects`)
  - ...

# MongoDB 3.2 Limitations

- Generic error message
  - Application needs to figure out what part of the constraints failed
- Cannot compare 1 key with another
  - Either within the same document or between documents
- Some operations not supported:
  - `$text`, `$geoNear`, `$near`, `$nearSphere`, `$where`
- Applications responsibility to bring legacy data into compliance with new rules
  - No audit or tools

# What validations remain in the app

- User interface
  - Don't have the database be the first place to detect that an email is poorly formatted
- Any validations that involve comparisons with
  - Other data in the **same document**
  - Data from **other documents**
  - **External information** (e.g. time of day)
- Semantic checks that are designed to fail frequently
  - e.g. user is in wrong country to use this service
  - Database should typically be testing for coding errors rather than implementing your business logic
- Determining why the database rejected a document in order to provide a meaningful error to the user



# Where MongoDB Validation excels<sub>(vs. RDBMS)</sub>

- Simple
  - Use familiar search expressions (MQL)
  - No need for stored procedures
- Flexible
  - Only enforced on mandatory parts of the schema
  - Can start adding new data at any point and then add validation later **if needed**
- Practical to deploy
  - Simple to roll out new rules across thousands of production servers
- Light weight
  - Negligible impact to performance

# Cleaning up legacy data

- Validator does not check if **existing** documents in the collection meet the new validation rules
- User/app can execute a query to identify & update any document which don't meet the new rules
  - Use **\$nor** on the full expression
- Be cautious about system impacts:
  - Could push working data set out of memory
  - Extra load if many documents need to be updated
  - Execute on secondary

```
secondary> db.runCommand({collMod: "bleh",  
                           validator: {  
                               a: {$lt:4}  
                           }  
                       });
```

```
secondary> db.bleh.find(  
                a: {$not: {$lt:4}}  
            ).count()
```

```
secondary> db.bleh.update(  
                {a: {$not: {$lt:4}}},  
                {$set: {a:3}},  
                {multi:true})
```

# Controlling validation

		validationLevel		
		off	moderate	strict
validationAction	warn	No checks	Warn on validation failure for inserts & updates to existing valid documents. Updates to existing invalid docs OK.	Warn on any validation failure for any insert or update.
	error	No checks	Reject invalid inserts & updates to existing valid documents. Updates to existing invalid docs OK.	Reject any violation of validation rules for any insert or update. <b>DEFAULT</b>

# Controlling validation

- Set behavior:

```
db.bleh.runCommand("collMod",  
                    {validationLevel: "moderate",  
                     validationAction: "warn"})
```

- Note that the warnings are written to the log

# Lifecycle

## Hacking (Day one)

- No document validation
- Release quick & often



## Analyze De facto Schema

- MongoDB Compass



## Add document validation rules

- Query & fix existing docs
- Log any new documents that break rules:  

```
{validationLevel: "moderate",  
validationAction: "warn"}
```



## Fix application

- Follow all rules
- When no new problems being logged, enter strict mode:  

```
{validationLevel: "strict",  
validationAction: "error"}
```

Application uses  
new data  
(Application  
evolves/additional  
app)

- If not mandated, stop  
here



## Analyze De-facto Schema

- MongoDB Compass



## Add document validation rules

- Query & fix existing docs
- Log any new documents that break rules:  

```
{validationLevel: "moderate",  
validationAction: "warn"}
```



## Fix application

- Follow all rules
- When no new problems being logged, enter strict mode:  

```
{validationLevel: "strict",  
validationAction: "error"}
```

# Versioning of Validations (optional)

```
db.runCommand({
  collMod: "contacts",
  validator:
    { $or: [{ version: { "$exists": false } },
            { version: 1,
              $and: [{ Name: { "$exists": true } } ]
            },
            { version: 2,
              $and: [{ Name: { "$type": "string" } } ]
            }
          ]
    }
})
```

- Application can lazily update documents with an older version or with no version set at all

# **Step through selecting and deploying a simple Document Validation Rule**

filter sources

- clusterdb.basket
- clusterdb.bleh
- clusterdb.homePriceAnnually
- clusterdb.homeSales
- clusterdb.hottestLocations
- clusterdb.inventory
- clusterdb.orders**
- clusterdb.places
- clusterdb.postcodes
- clusterdb.products
- clusterdb.purchases
- clusterdb.sales
- clusterdb.stock
- flights.flightstats
- moreStuff.myCollection

This report is based on a sample of 100 documents:

# orders

DOCUMENTS

3.3k

total size  
206.5KB

avg. size  
65B

INDEXES

2

total size  
64.0KB

avg. size  
32.0KB

{ }

APPLY

RESET

\_id

number



2800 2257 2333 3398 1893 1424 3866 2854 2825  
3780 1423 3534 3168 3562 2967

item

string

null unc

abc jkl

price

number

string



quantity

number





# Drill down into anomalies

---

price

number

string

free

if you have to ask....

---

```
{"price":"if you have to ask...."}
```

APPLY

RESET

**\_id**

number



4214

4117

4078

4053

4052

4105

4079

4103

4152

4188

4112

4071

4150

**item**

string

abc

**price**

string

if you have to ask....

```
_id: 4086  
item: "abc"  
price: "if you have to ask...."  
quantity: 4
```

```
_id: 4108  
item: "abc"  
price: "if you have to ask...."  
quantity: 1
```

```
_id: 4241  
item: "abc"  
price: "if you have to ask...."  
quantity: 6
```

```
_id: 4156  
item: "abc"  
price: "if you have to ask...."  
quantity: 9
```

```
_id: 4142  
item: "abc"  
price: "if you have to ask...."  
quantity: 9
```

# 1. Prevent New Malformed Documents

```
> db.orders.runCommand("collMod",
    {validationLevel: "moderate",
     validationAction: "error"});

> db.runCommand({collMod: "orders",
    validator: {
      $or: [
        {price: {$type: 1}},
        {price: {$type: 16}},
        {price: {$type: 18}}
      ]
    }
  });

> db.getCollectionInfos({name:"orders"})
```

```
{
  "name": "orders",
  "options": {
    "validator": {
      "$or": [
        { "price": { "$type": 1 } },
        { "price": { "$type": 16 } },
        { "price": { "$type": 18 } }
      ]
    },
    "validationLevel": "moderate",
    "validationAction": "error"
  }
}
```

## 2. Prevent New Malformed Documents

```
> db.orders.insert({  
  "_id": 6666,  
  "item": "jkl",  
  "price": "rogue",  
  "quantity": 1 });
```

```
Document failed validation  
WriteResult({  
  "nInserted": 0,  
  "writeError": {  
    "code": 121,  
    "errmsg": "Document failed validation"  
  }  
})
```

### 3. Clean-Up Legacy Documents

```
> db.orders.findOne(  
  {price:  
    {$type: "string"}}  
);
```

```
> db.orders.update(  
  {_id: 3500},  
  {$set:{quantity: 12}}  
);
```

```
{  
  "_id": 3500,  
  "item": "abc",  
  "price": "free",  
  "quantity": 8  
}
```

Updated 1 existing record(s) in 6ms

```
WriteResult(  
  "nMatched": 1,  
  "nUpserted": 0,  
  "nModified": 1  
)
```

# 3. Clean-Up Legacy Documents

```
> db.orders.update(  
  {price:"free"},  
  {$set: {price: 0}},  
  {multi: true});
```

```
> db.orders.update(  
  {price:"if you have to ask...."},  
  {$set: {price: 1000000}},  
  {multi: true});
```

## 4. Confirm Results



```
> db.orders.find(  
  {$nor: [  
    {price: {$type: 1}},  
    {price: {$type: 16}},  
    {price: {$type: 18}}  
  ]})
```

Fetches 0 record(s) in **5ms**

# Next Steps

- Document Validation - Adding Just the Right Amount of Control Over Your Documents
  - <https://www.mongodb.com/blog/post/document-validation-part-1-adding-just-the-right-amount-of-control-over-your-documents>
- “Document Validation and What Dynamic Schema Means” – Eliot Horowitz
  - <http://www.eliothorowitz.com/blog/2015/09/11/document-validation-and-what-dynamic-schema-means/>
- “Bulletproof Data Management” – MongoDB World 2015
  - <https://www.mongodb.com/presentations/data-management-3-bulletproof-data-management>
- Documentation
  - <http://docs.mongodb.org/manual/release-notes/3.1-dev-series/#document-validation>
- Not yet ready for production but download and try MongoDB 3.2 RC
  - <https://www.mongodb.org/downloads#development>
- Feedback
  - <https://www.mongodb.com/blog/post/announcing-the-mongodb-3-2-bug-hunt>
  - <https://jira.mongodb.org/>

DISCLAIMER: MongoDB's product plans are for informational purposes only. MongoDB's plans may change and you should not rely on them for delivery of a specific feature at a specific time.