

mongoDB

Distributed Consensus in MongoDB

Spencer T Brody

Senior Software Engineer at MongoDB

@stbrody

Agenda

- Introduction to consensus
- Leader-based replicated state machine
- Elections and data replication in MongoDB
- Improvements coming in MongoDB 3.2

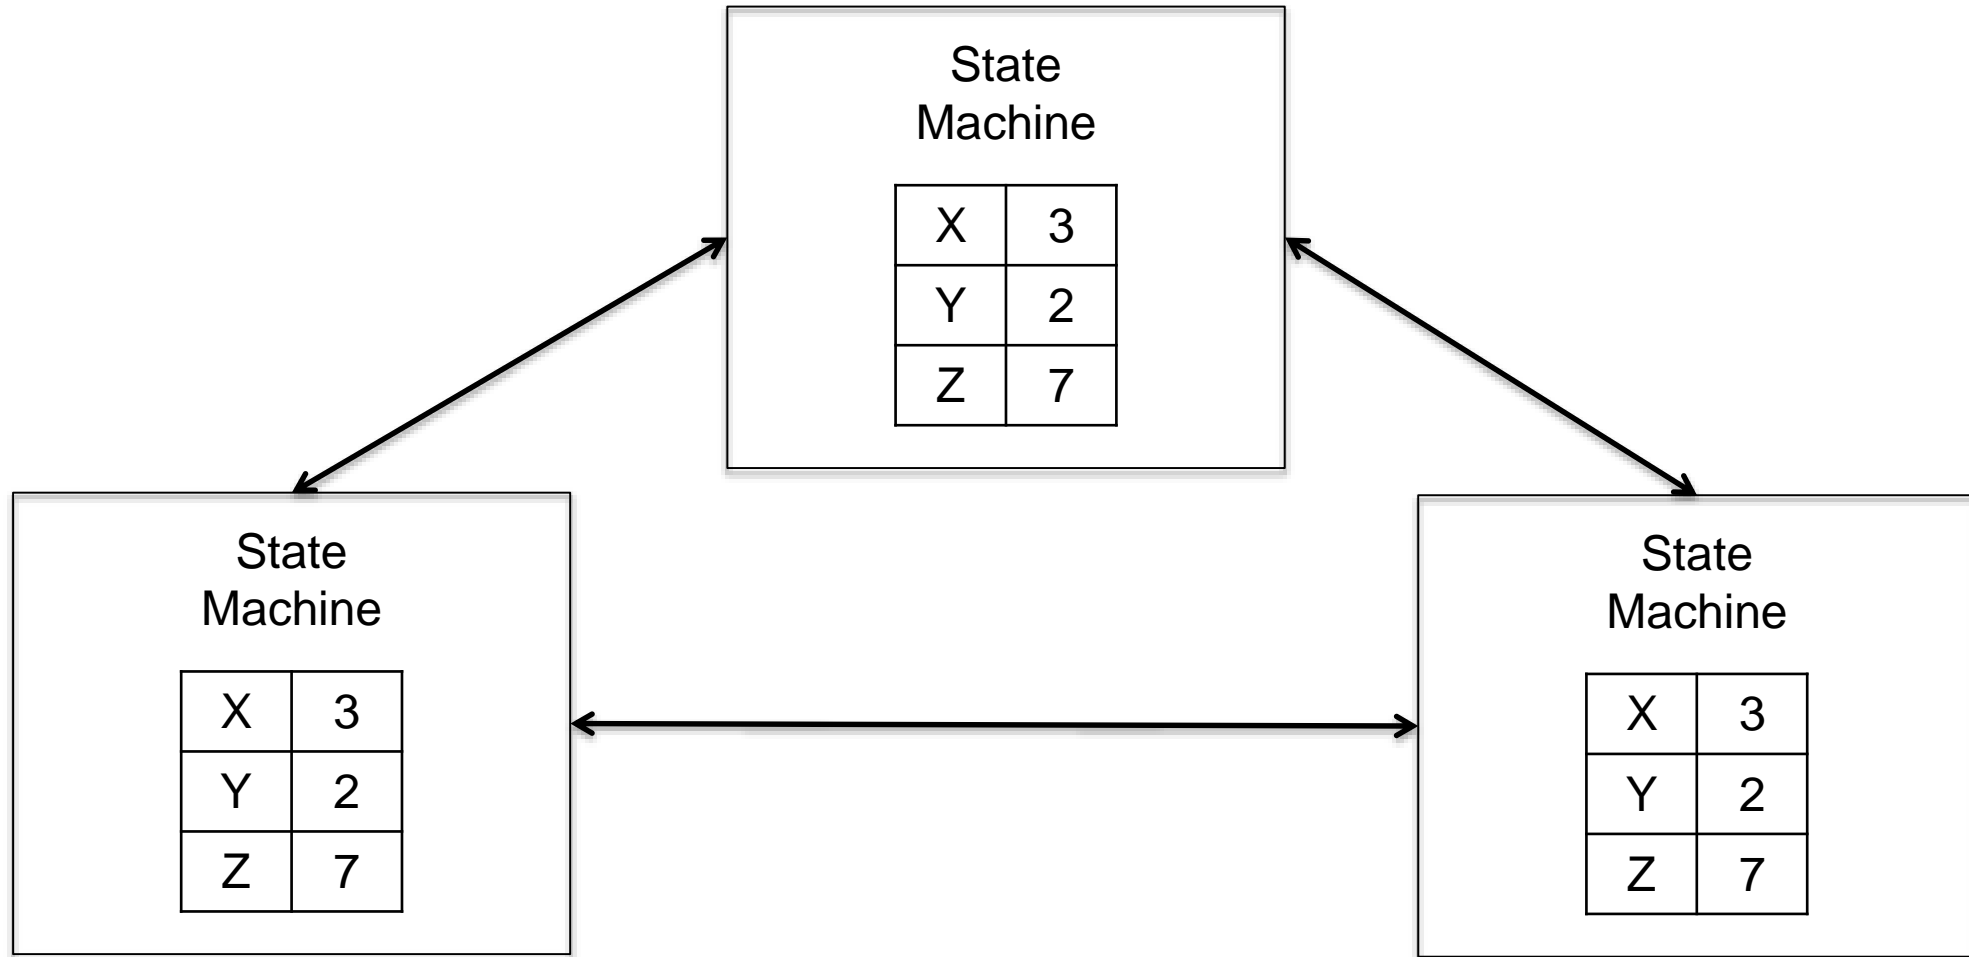
Why use Replication?

- Data redundancy
- High availability

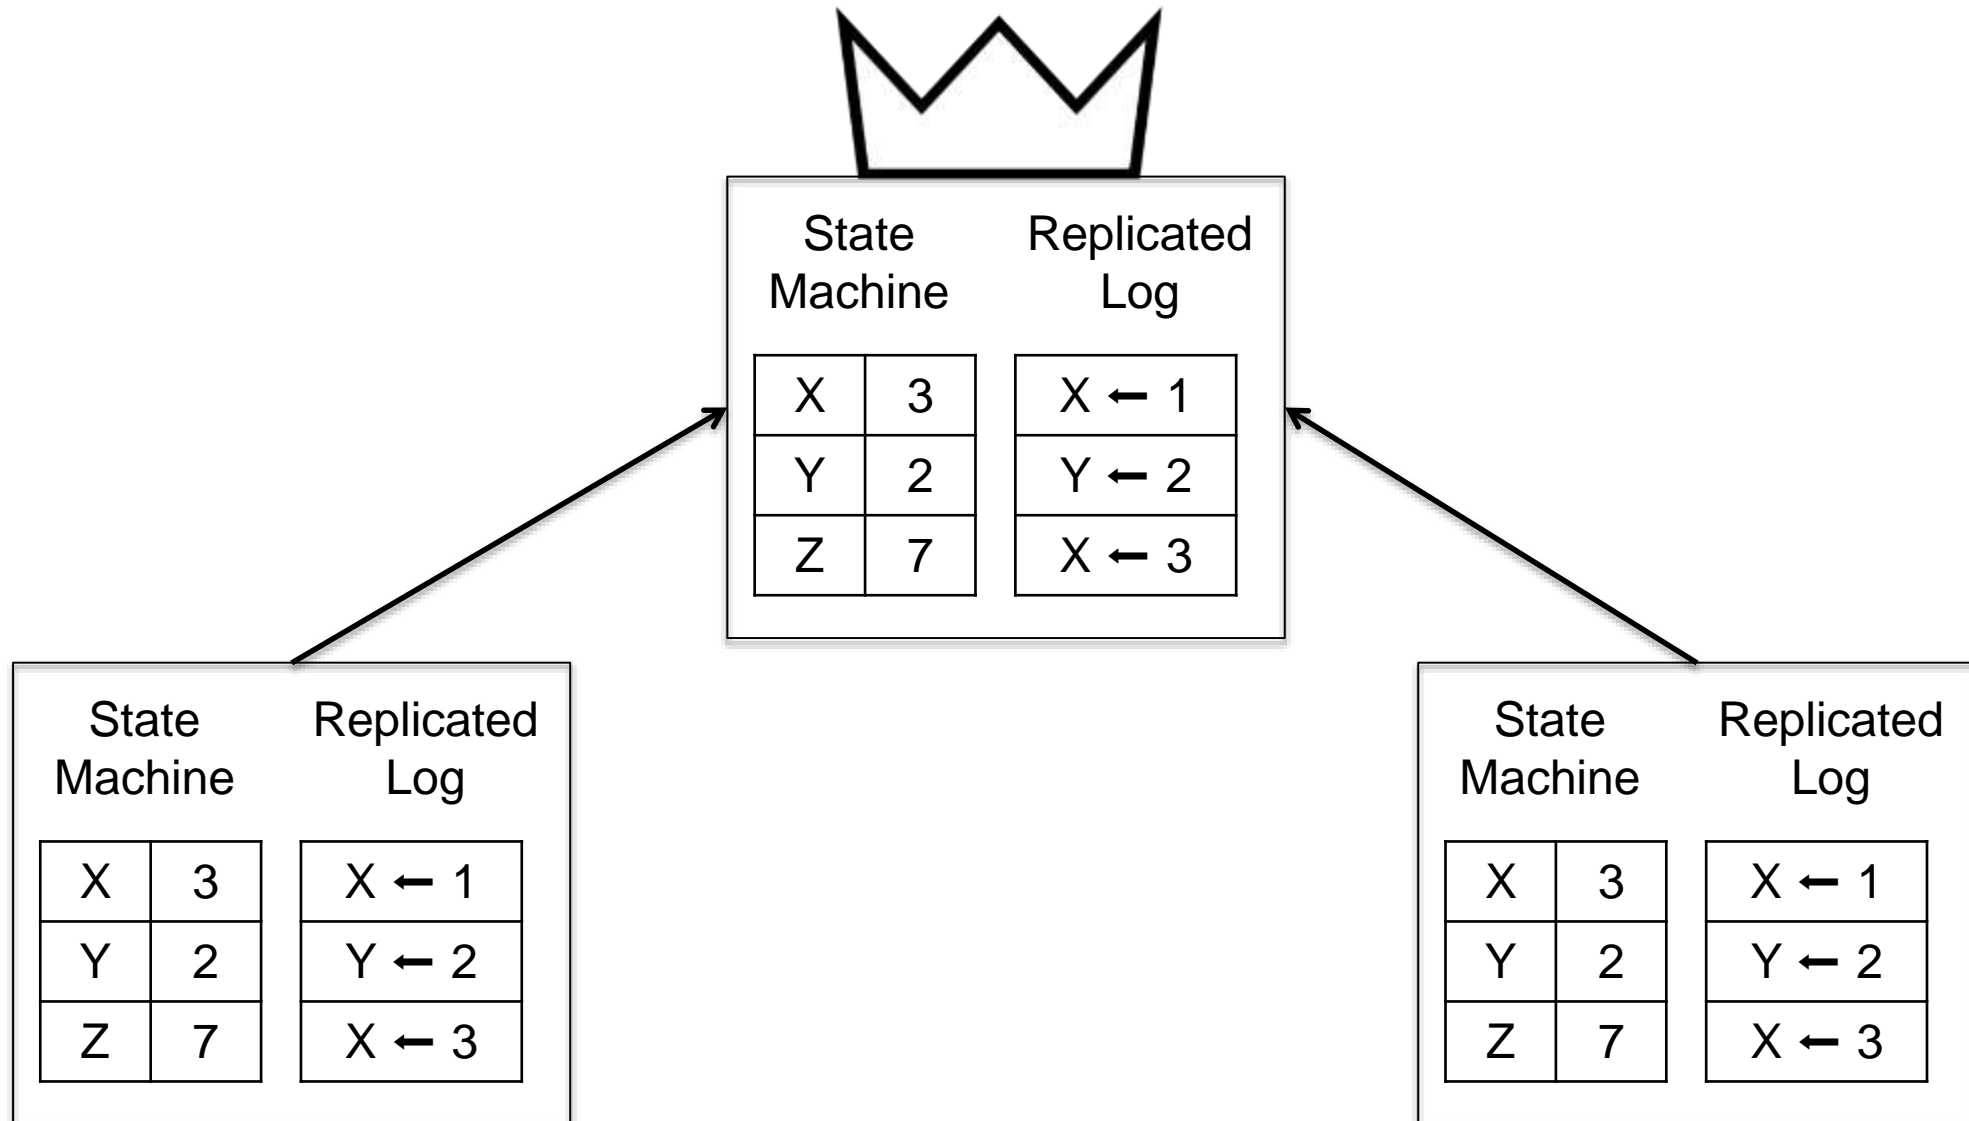
What is Consensus?

- Getting multiple processes/servers to agree on something
- Must handle a wide range of failure modes
 - Disk failure
 - Network partitions
 - Machine freezes
 - Clock skews

Basic consensus



Leader Based Consensus



Agenda

- Introduction to consensus
- Leader-based replicated state machine
- Elections and data replication in MongoDB
- Improvements coming in MongoDB 3.2

Elections



X	3	X ← 1
Y	2	Y ← 2
Z	7	X ← 3

X	3	X ← 1
Y	2	Y ← 2
Z	7	X ← 3

X	3	X ← 1
Y	2	Y ← 2
Z	7	X ← 3

X	3	X ← 1
Y	2	Y ← 2
Z	7	X ← 3

X	3	X ← 1
Y	2	Y ← 2
Z	7	X ← 3

Elections

X	3	X ← 1
Y	2	Y ← 2
Z	7	X ← 3

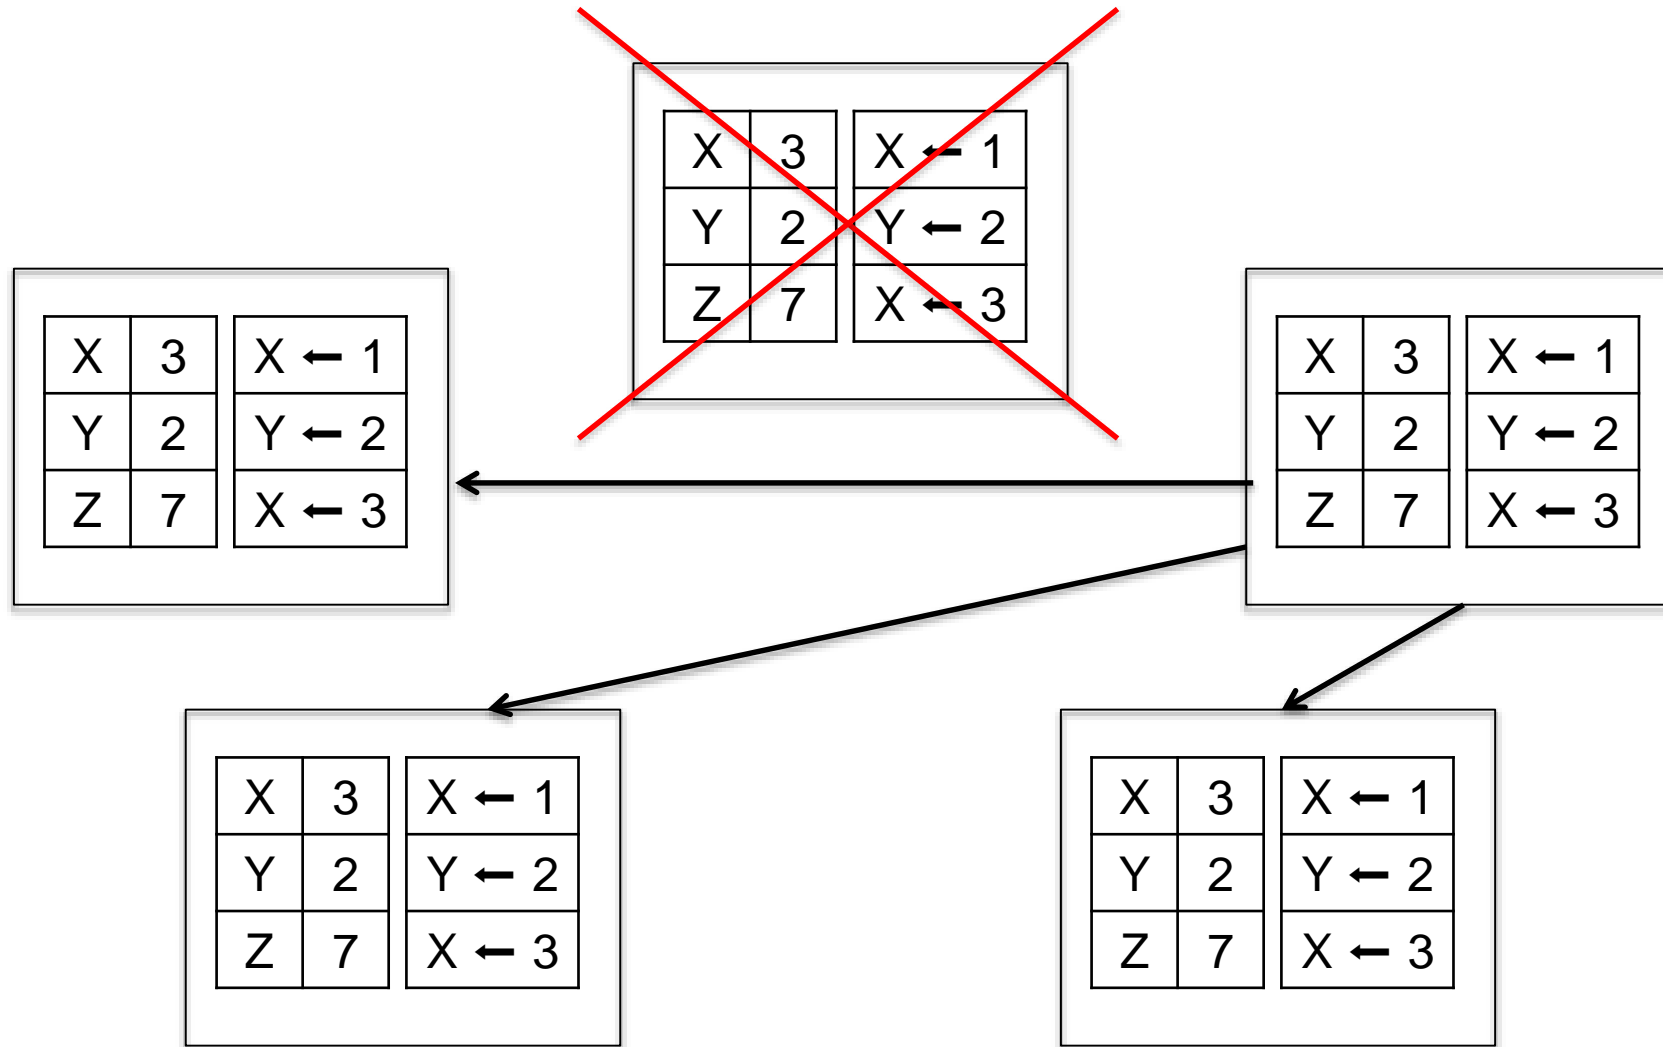
X	3	X ← 1
Y	2	Y ← 2
Z	7	X ← 3

X	3	X ← 1
Y	2	Y ← 2
Z	7	X ← 3

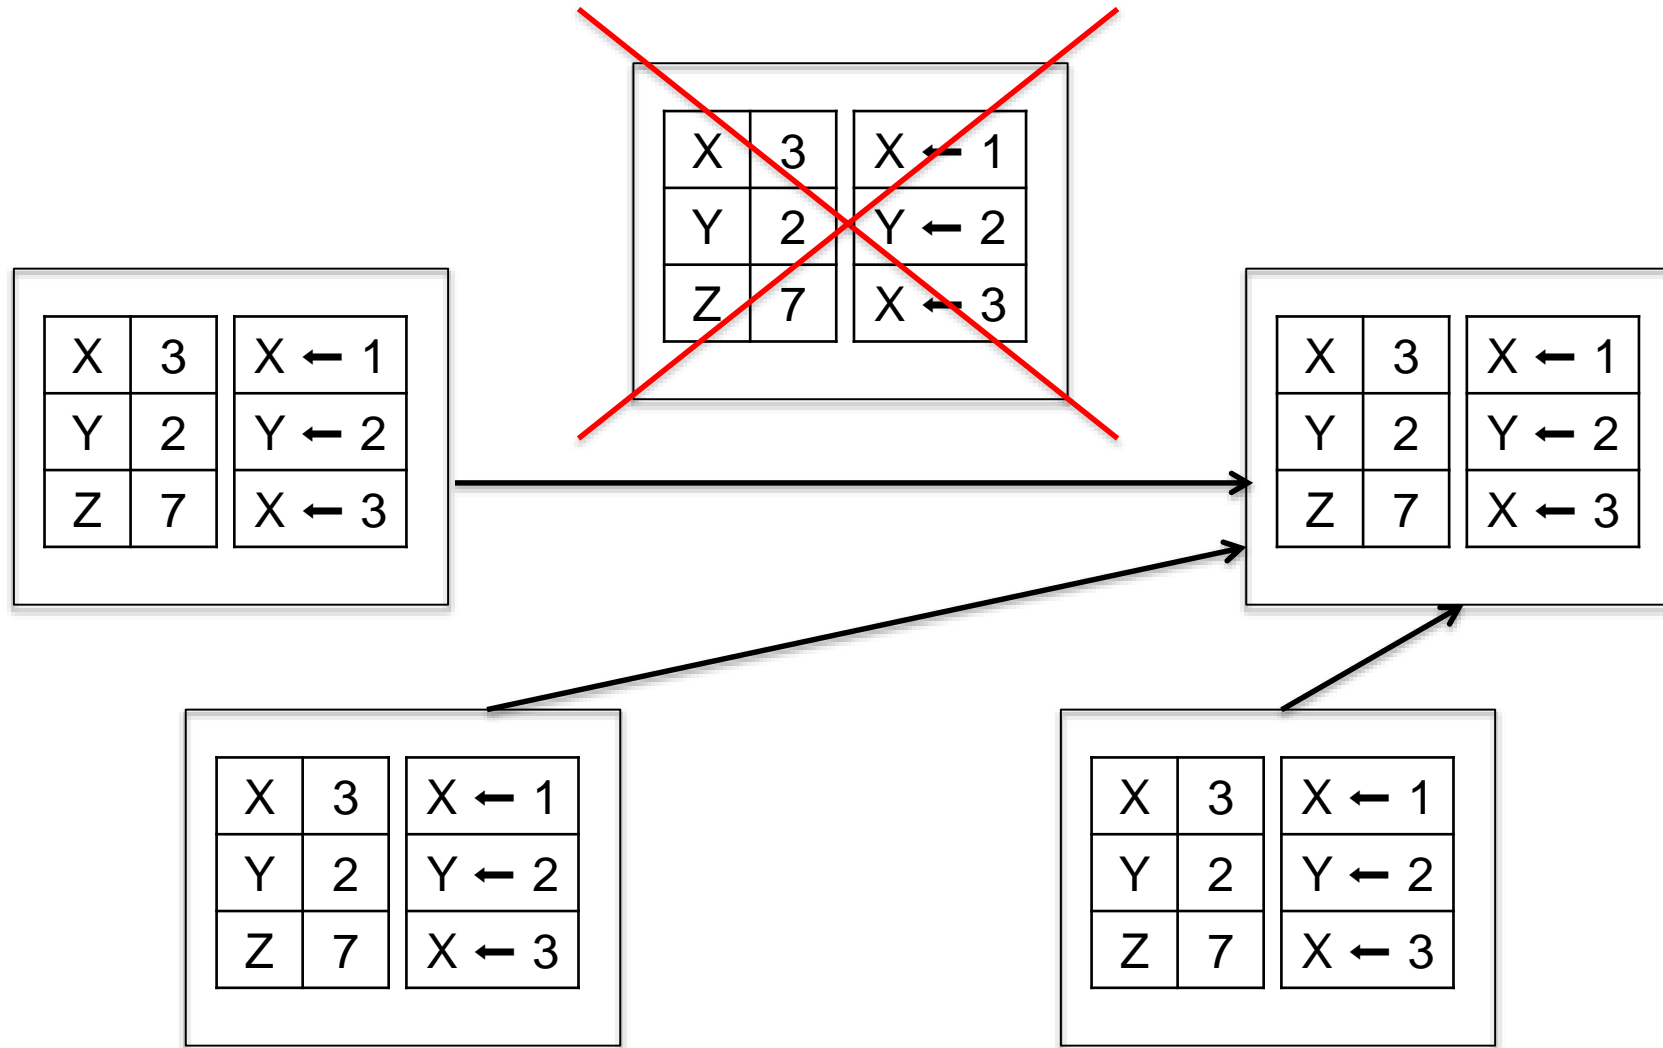
X	3	X ← 1
Y	2	Y ← 2
Z	7	X ← 3

X	3	X ← 1
Y	2	Y ← 2
Z	7	X ← 3

Elections



Elections




Elections

X	3	X ← 1
Y	2	Y ← 2
Z	7	X ← 3

X	3	X ← 1
Y	2	Y ← 2
Z	7	X ← 3

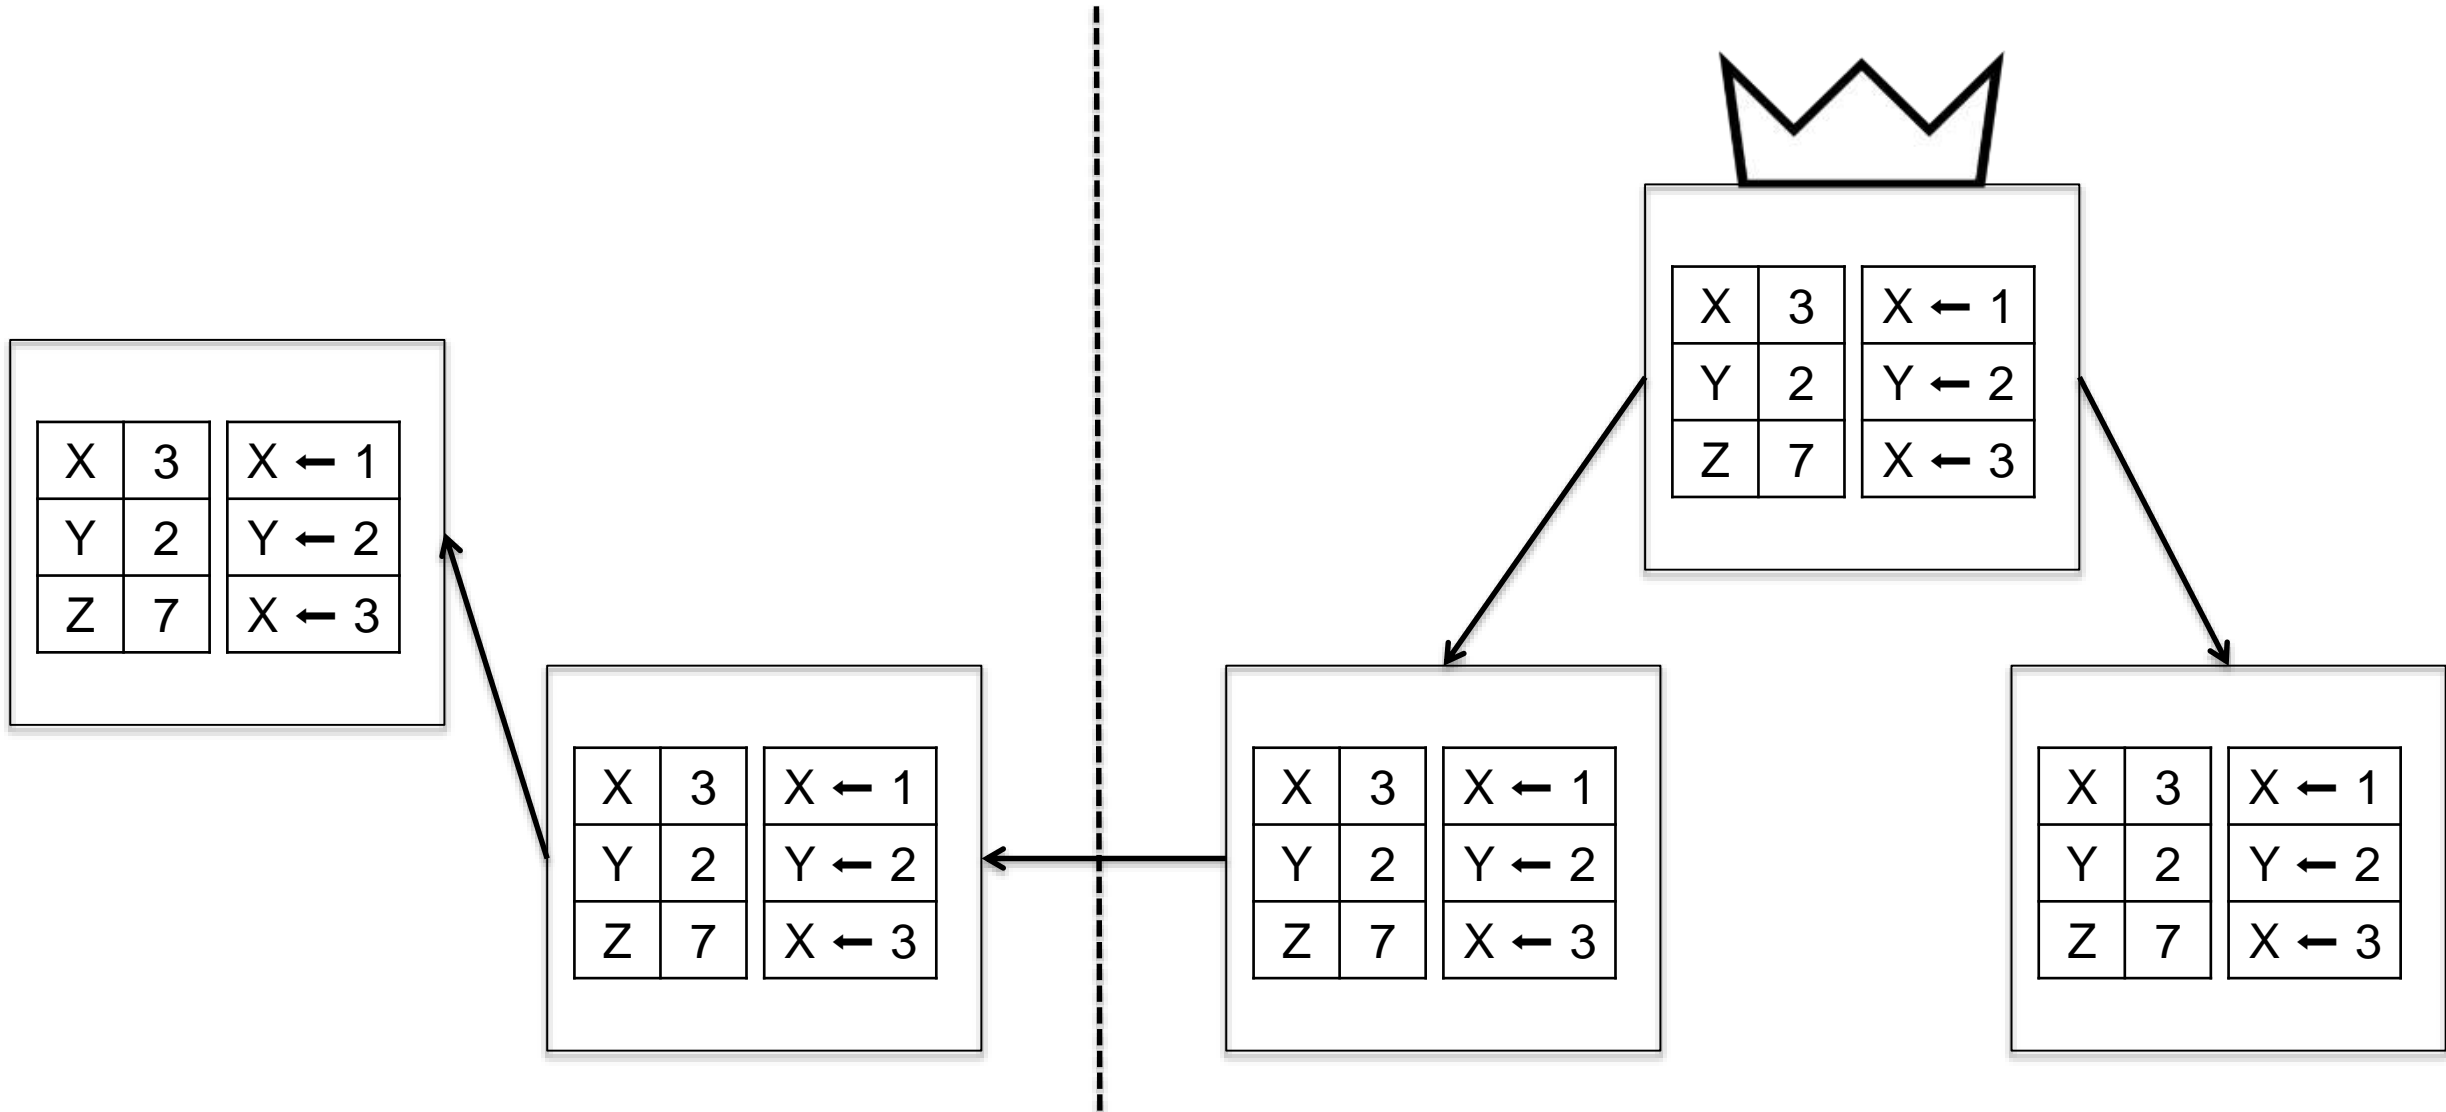
X	3	X ← 1
Y	2	Y ← 2
Z	7	X ← 3



X	3	X ← 1
Y	2	Y ← 2
Z	7	X ← 3

X	3	X ← 1
Y	2	Y ← 2
Z	7	X ← 3

Data Replication



Agenda

- Introduction to consensus
- Leader-based replicated state machine
- Elections and data replication in MongoDB
- Improvements coming in MongoDB 3.2

Agenda

- Introduction to consensus
- Leader-based replicated state machine
- Elections and data replication in MongoDB
- Improvements coming in MongoDB 3.2
 - Goals and inspiration from Raft Consensus Algorithm
 - Preventing double voting
 - Monitoring node status
 - Calling for elections

Goals for MongoDB 3.2

- Decrease failover time
- Speed up detection and resolution of false primary situations

Finding Inspiration in Raft

- “In Search of an Understandable Consensus Algorithm” by Diego Ongaro: <https://ramcloud.stanford.edu/raft.pdf>
- Designed to address the shortcomings of Paxos
 - Easier to understand
 - Easier to implement in real applications
- Provably correct
- Remarkably similar to what we’re doing already

Raft Concepts

- Term (election) IDs
- Monitoring node status using existing data replication channel
- Asymmetric election timeouts

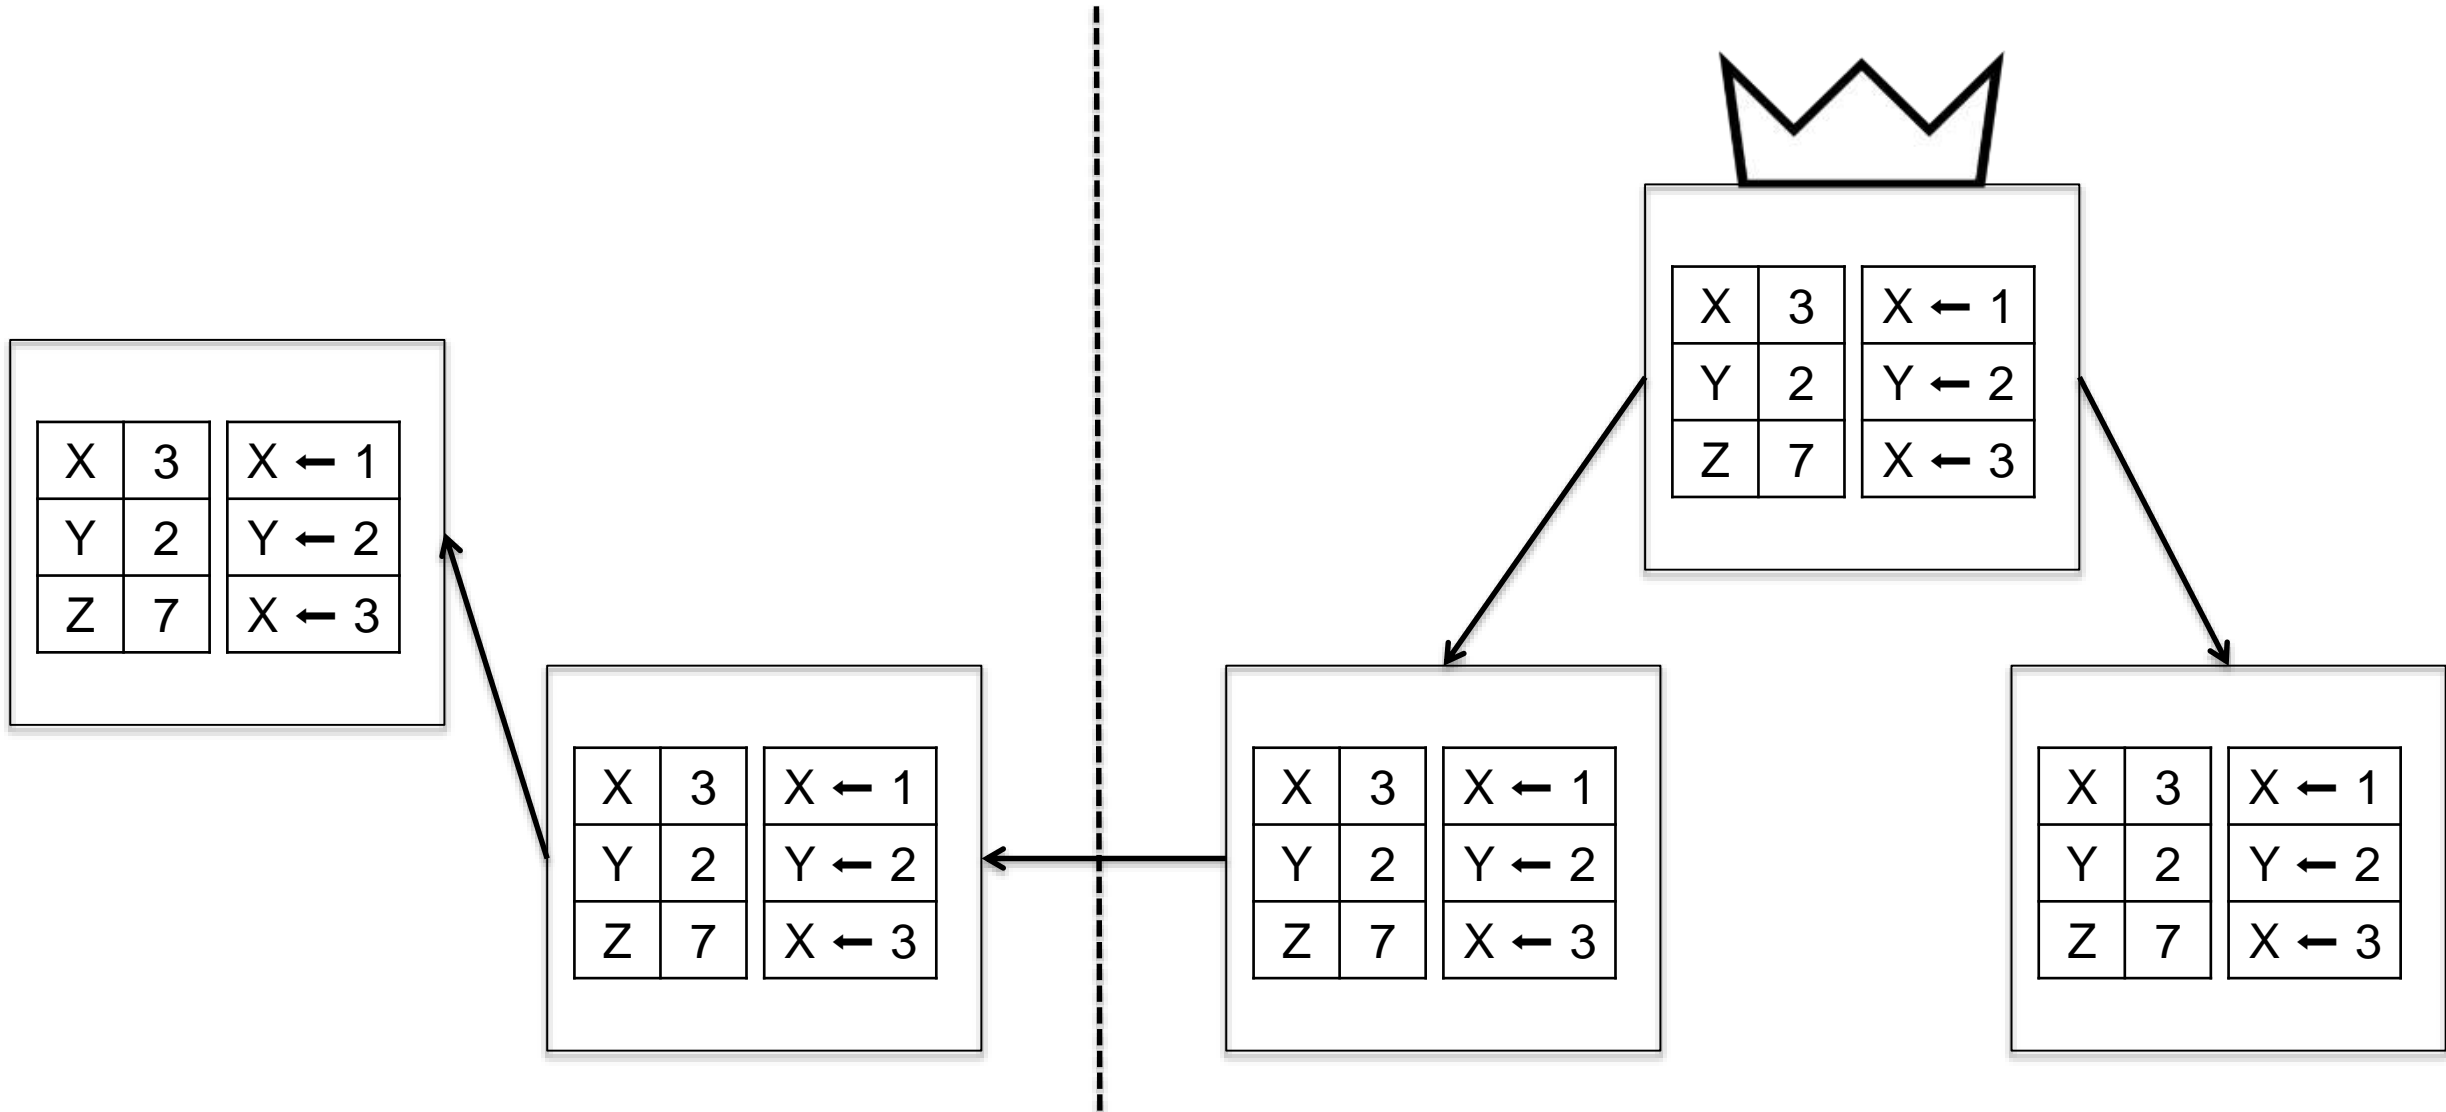
Preventing Double Voting

- Can't vote for 2 nodes in the same election
- Pre-3.2: 30 second vote timeout
- Post-3.2: Term IDs
- Term:
 - Monotonically increasing ID
 - Incremented on every election *attempt*
 - Lets voters distinguish elections so they can vote twice quickly in different elections.

Monitoring Node Status

- Pre-3.2: Heartbeats
 - Sent every two seconds from every node to every other node
 - Volume increases quadratically as nodes are added to the replica set
- Post-3.2: Extra metadata sent via existing data replication channel
 - Utilizes chained replication
 - Faster heartbeats = faster elections and detection of false primaries

Data Replication



Determining When To Call For An Election

- Tradeoff between failover time and spurious failovers
- Node calls for an election when it hasn't heard from the primary within the election timeout
- Starting in 3.2:
 - Election timeout is configurable
 - Election timeout is varied randomly for each node
 - Varying timeouts help reduce tied votes
 - Fewer tied votes = faster failover

Conclusion

- MongoDB 3.2 will have
 - Faster failovers
 - Faster error detection
 - More control to prevent spurious failovers
- This means your systems are
 - More stable
 - More resilient to failure
 - Easier to maintain

Questions?